

第二周作业

1. 基于dockerfile, 实现分层构建的nginx业务镜像
2. 基于docker实现对容器的CPU和内存的资源限制
3. 部署http协议的harbor镜像仓库
4. 基于docker-compose实现对nginx+tomcat web服务的单机编排
5. 扩展作业:

5.1 掌握containerd的安装

5.2 基于nerdctl拉取镜像和创建容器

1. 基于dockerfile, 实现分层构建的nginx业务镜像

```
1 #清除历史容器
2 root@castillo:~# docker stop `docker ps | awk '{if (NR>1){print $1}}'`
3 root@castillo:~# docker rm `docker ps -a | awk '{if (NR>1){print $1}}'`
```

1.1. 编写二进制安装nginxdockerfile

```
1 FROM ubuntu:22.04
2 MAINTAINER "jinliucastillo"
3
4 ADD sources.list /etc/apt/sources.list
5
6 RUN apt update && apt install -y iproute2 ntpdate tcpdump telnet traceroute nfs-
  kernel-server nfs-common lrzsz tree openssl libssl-dev libpcre3 libpcre3-dev zlib1g-
  dev ntpdate tcpdump telnet traceroute gcc openssh-server lrzsz tree openssl libssl-
  dev libpcre3 libpcre3-dev zlib1g-dev ntpdate tcpdump telnet traceroute iotop unzip
  zip make
7
8 ADD nginx-1.22.1.tar.gz /usr/local/src/
9 RUN cd /usr/local/src/nginx-1.22.1 && ./configure --prefix=/apps/nginx && make &&
  make install && ln -sv /apps/nginx/sbin/nginx /usr/bin
10 RUN groupadd -g 2088 nginx && useradd -g nginx -s /usr/sbin/nologin -u 2088 nginx &&
  chown -R nginx.nginx /apps/nginx
11 ADD nginx.conf /apps/nginx/conf/
12 ADD frontend.tar.gz /apps/nginx/html/
13
14 EXPOSE 80 443
15 CMD ["nginx","-g","daemon off;"]
```

1.2. 构建nginx镜像

```

1 docker build -t harbor.jinliu.net/myserver/nginx:v1 .
2 root@castillo:/opt/cicd# docker images
3 REPOSITORY          TAG          IMAGE ID          CREATED          SIZE
4 harbor.jinliu.net/myserver/nginx    v1          5f227fbc596b     About a minute ago    470MB
5 ubuntu              22.04       9d28ccdc1fc7     11 months ago       76.3MB
6 root@castillo:/opt/cicd# docker run -dit -p 80:80 harbor.jinliu.net/myserver/nginx:v1

```

1.3. 提交镜像

```

1 root@castillo:/opt/cicd# docker commit -m "nginx image v2" -a "jinliu castillo" -c
"EXPOSE 80 443" 2330728e0b6e harbor.jinliu.net/myserver/nginx:v2
2 sha256:e6954854a75e849fc0cb4164da30f2b5dc52d8ac58344e76014eacaf7de2ef3f
3 root@castillo:/opt/cicd# docker images
4 REPOSITORY          TAG          IMAGE ID          CREATED          SIZE
5 harbor.jinliu.net/myserver/nginx    v2          e6954854a75e     6 seconds ago     514MB
6 harbor.jinliu.net/myserver/nginx    v1          5f227fbc596b     21 minutes ago    470MB
7 ubuntu              22.04       9d28ccdc1fc7     11 months ago     76.3MB
8 root@castillo:/opt/cicd#

```

2. 基于docker实现对容器的CPU和内存的资源限制

◆对于Linux 主机，如果没有足够的内容来执行其他重要的系统任务，将会抛出 OOM (Out of Memory Exception,内存溢出、内存泄漏、内存异常)，随后系统会开始杀死进程以释放内存，凡是运行在宿主机的进程都有可能被 kill，包括Dockerd和其它的应用程序，如果重要的系统进程被Kill,会导致和该进程相关的服务全部宕机。

```

Oct 10 17:09:28 docker-server1 kernel: [57402.792264] oom-kill:constraint=CONSTRAINT_NONE,nodemask=(null),cpuset=systemd-
journal.service,mems_allowed=0,global_om,task_memcg=/system.slice/docker-81aca8d9f526f25eea0015af00271475c9e5a794b762fa70dfe7af2f7efaff64.scope,task=stress-ng-
vm,pid=41298,uid=0
Oct 10 17:09:28 docker-server1 kernel: [57402.792303] Out of memory: Killed process 41298 (stress-ng-vm) total-vm:530572kB, anon-rss:314588kB,
file-rss:4kB, shme
m-rss:24kB, UID:0 pgtables:668kB oom_score_adj:1000
Oct 10 17:09:29 docker-server1 kernel: [57403.278032] sshd invoked oom-killer: gfp_mask=0x1100cca(GFP_HIGHUSER_MOVABLE), order=0,
oom_score_adj=0
Oct 10 17:09:29 docker-server1 kernel: [57403.278038] CPU: 0 PID: 39756 Comm: sshd Not tainted 5.15.0-43-generic #46-Ubuntu
Oct 10 17:09:29 docker-server1 kernel: [57403.278040] Hardware name: VMware, Inc. VMware Virtual Platform/440BX Desktop Reference Platform, BIOS
6.00 07/22/2020
Oct 10 17:09:29 docker-server1 kernel: [57403.278041] Call Trace:
Oct 10 17:09:29 docker-server1 kernel: [57403.278042] <TASK>
Oct 10 17:09:29 docker-server1 kernel: [57403.278044] show_stack+0x52/0x58
Oct 10 17:09:29 docker-server1 kernel: [57403.278048] dump_stack_lvl+0x4a/0x5f
Oct 10 17:09:29 docker-server1 kernel: [57403.278050] dump_stack+0x10/0x12
Oct 10 17:09:29 docker-server1 kernel: [57403.278052] dump_header+0x53/0x224
Oct 10 17:09:29 docker-server1 kernel: [57403.278054] oom_kill_process.cold+0xb/0x10
Oct 10 17:09:29 docker-server1 kernel: [57403.278055] out_of_memory+0x106/0x2e0
Oct 10 17:09:29 docker-server1 kernel: [57403.278058] __alloc_pages_slowpath.constprop.0+0x97a/0xa40

```

```

1 #查看内核支持情况
2 root@castillo:/opt/cicd# docker info
3 Client:
4 Context:    default
5 Debug Mode: false
6 Plugins:
7   app: Docker App (Docker Inc., v0.9.1-beta3)
8   buildx: Docker Buildx (Docker Inc., v0.9.1-docker)
9   compose: Docker Compose (Docker Inc., v2.10.2)
10  scan: Docker Scan (Docker Inc., v0.17.0)
11
12 Server:
13 Containers: 4

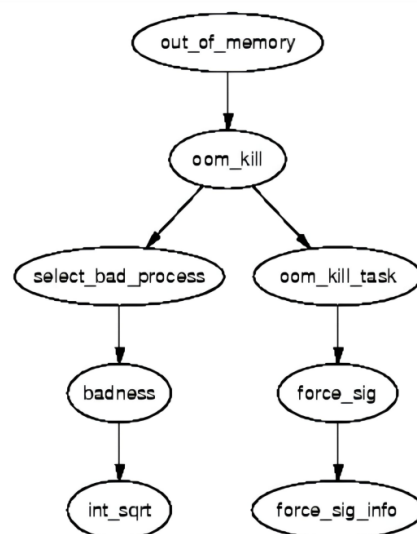
```

```
14   Running: 1
15   Paused: 0
16   Stopped: 3
17   Images: 12
18   Server Version: 20.10.18
19   Storage Driver: overlay2
20     Backing Filesystem: extfs
21     Supports d_type: true
22     Native Overlay Diff: true
23     userxattr: false
24   Logging Driver: json-file
25   Cgroup Driver: systemd
26   Cgroup Version: 2
27   Plugins:
28     Volume: local
29     Network: bridge host ipvlan macvlan null overlay
30     Log: awslogs fluentd gcplogs gelf journald json-file local logentries splunk syslog
31   Swarm: inactive
32   Runtimes: io.containerd.runc.v2 io.containerd.runtime.v1.linux runc
33   Default Runtime: runc
34   Init Binary: docker-init
35   containerd version: 9cd3357b7fd7218e4aec3eae239db1f68a5a6ec6
36   runc version: v1.1.4-0-g5fd4c4d
37   init version: de40ad0
38   Security Options:
39     apparmor
40     seccomp
41       Profile: default
42     cgroupns
43   Kernel Version: 5.15.0-52-generic
44   Operating System: Ubuntu 22.04.1 LTS
45   OSType: linux
46   Architecture: x86_64
47   CPUs: 2
48   Total Memory: 3.832GiB
49   Name: castillo
50   ID: YUYJ:XCZZ:VP76:ENLD:AFBF:PUR7:2WMK:ENSP:VMYB:WHWT:JKCL:MFNL
51   Docker Root Dir: /var/lib/docker
52   Debug Mode: false
53   Registry: https://index.docker.io/v1/
54   Labels:
55     Experimental: false
56   Insecure Registries:
57     127.0.0.0/8
58   Registry Mirrors:
59     https://rlyvli78.mirror.aliyuncs.com/
60   Live Restore Enabled: false
61
62   root@castillo:/opt/cicd#
```

2.1. oom优先级机制

◆OOM优先级机制:

- linux会为每个进程算一个分数，最终他会将分数最高的进程kill。
- `/proc/PID/oom_score_adj` #范围为-1000到1000，值越高越容易被宿主机kill掉，如果将该值设置为-1000，则进程永远不会被宿主机kernel kill。
- `/proc/PID/oom_adj` #范围为-17到+15，取值越高越容易被干掉，如果是-17，则表示不能被kill，该设置参数的存在是为了和旧版本的Linux内核兼容。
- `/proc/PID/oom_score` #这个值是系统综合进程的内存消耗量、CPU时间(utime + stime)、存活时间(uptime - start time)和oom_adj计算出的进程得分，消耗内存越多得分越高，越容易被宿主机kernel强制杀死。



◆物理内存限制参数:

- `-m` or `--memory` #限制容器可以使用的最大内存量，如果设置此选项，最小存值为4m（4兆字节）。
- `--memory-swap` #容器可以使用的交换分区大小，必须要在设置了物理内存限制的前提才能设置交换分区的限制
- `--memory-swappiness` #设置容器使用交换分区的倾向性，值越高表示越倾向于使用swap分区，范围为0-100，0为能不用就不用，100为能用就用。
- `--kernel-memory` #容器可以使用的最大内核内存量，最小为4m，由于内核内存与用户空间内存隔离，因此无法与用户空间内存直接交换，因此内核内存不足的容器可能会阻塞宿主机资源，这会对主机和其他容器或者其他服务进程产生影响，因此不要设置内核内存大小。
- `--memory-reservation` #允许指定小于`--memory`的软限制，当Docker检测到主机上的争用或内存不足时会激活该限制，如果使用`--memory-reservation`，则必须将其设置为低于`--memory`才能使其优先。因为它是软限制，所以不能保证容器不超过限制。
- `--oom-kill-disable` #默认情况下，发生OOM时，kernel会杀死容器内进程，但是可以使用`--oom-kill-disable`参数，可以禁止oom发生在指定的容器上，即 仅在已设置`-m` / `--memory`选项的容器上禁用OOM，如果`-m` 参数未配置，产生OOM时，主机为了释放内存还会杀死系统进程。

2.2. 测试内存限制

```
1 root@castillo:/opt/cicd# docker run -it --rm --name jinliu-c1 lorel/docker-stress-ng
--vm 2 --vm-bytes 256M
2 stress-ng: info: [1] defaulting to a 86400 second run per stressor
3 stress-ng: info: [1] dispatching hogs: 2 vm
4
5 root@castillo:~# docker stats
6 CONTAINER ID   NAME          CPU %       MEM USAGE / LIMIT   MEM %       NET I/O
7 5dd8008128b8   jinliu-c1     198.06%    515.8MiB / 3.832GiB  13.15%      796B / 0B
8 2330728e0b6e   sleepy_johnson 0.00%      15.47MiB / 3.832GiB  0.39%      11MB /
110kB
9
10 root@castillo:~# docker run -it --rm -m 256m --name jinliu-c2 lorel/docker-stress-ng
--vm 2 --vm-bytes 256M
11 Given value bytes is not a valid decimal for the vm option
12 root@castillo:~
```

```

13 root@castillo:~# docker run -it --rm -m 256m --name jinliu-c2 lorel/docker-stress-ng
    --vm 2 --vm-bytes 256M
14 stress-ng: info: [1] defaulting to a 86400 second run per stressor
15 stress-ng: info: [1] dispatching hogs: 2 vm
16
17
18 root@castillo:~# docker stats
19 CONTAINER ID   NAME          CPU %       MEM USAGE / LIMIT   MEM %      NET I/O
    BLOCK I/O      PIDS
20 2330728e0b6e   sleepy_johnson 0.00%       15.47MiB / 3.832GiB  0.39%      11MB /
    110kB  0B / 121MB      2
21 716e0f480798   jinliu-c2      17.54%      255.9MiB / 256MiB   99.95%     656B / 0B
    99.5MB / 3.86GB  5

```

2.3. 测试cpu限制

✓ 注：CPU资源限制是将分配给容器的2核心分配到了宿主机每一核心CPU上，也就是容器的总CPU值是在宿主机的每一个核心CPU分配了部分比例。

- Tasks: 288 total, 10 running, 278 sleeping, 0 stopped, 0 zombie
- %Cpu0 : 51.2 us, 0.0 sy, 0.0 ni, 48.8 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
- %Cpu1 : 26.4 us, 23.4 sy, 0.0 ni, 50.2 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
- %Cpu2 : 23.1 us, 27.7 sy, 0.0 ni, 49.2 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
- %Cpu3 : 18.3 us, 31.3 sy, 0.0 ni, 50.3 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st

```

1 root@castillo:~# docker run -it --rm --name jjinliu-c1 --cpus 2 lorel/docker-stress-ng
    --cpu 4 --vm 4
2 stress-ng: info: [1] defaulting to a 86400 second run per stressor
3 stress-ng: info: [1] dispatching hogs: 4 cpu, 4 vm
4 ^Cstress-ng: info: [1] successful run completed in 91.84s
5 root@castillo:~#

```

2.4. 将容器运行到指定的CPU上

```

1 root@castillo:~# docker run -it --rm --name jinliu-c1 --cpus 1 --cpuset-cpus 1
    lorel/docker-stress-ng --cpu 2 --vm 2
2 stress-ng: info: [1] defaulting to a 86400 second run per stressor
3 stress-ng: info: [1] dispatching hogs: 2 cpu, 2 vm
4
5 root@castillo:~# docker stats
6 CONTAINER ID   NAME          CPU %       MEM USAGE / LIMIT   MEM %      NET I/O
    BLOCK I/O      PIDS
7 2330728e0b6e   sleepy_johnson 0.00%       15.47MiB / 3.832GiB  0.39%      11MB /
    110kB  0B / 121MB      2
8 6fbae7fcc895   jinliu-c1      100.35%     519.9MiB / 3.832GiB  13.25%     726B / 0B
    0B / 0B        7
9
10 root@castillo:/opt/cicd# top
11 top - 17:08:28 up 10 days,  2:41,  8 users,  load average: 0.67, 1.56, 1.21
12 Tasks: 140 total,  5 running, 135 sleeping,  0 stopped,  0 zombie
13 %Cpu0  :  0.3 us,  0.3 sy,  0.0 ni, 99.0 id,  0.3 wa,  0.0 hi,  0.0 si,  0.0 st
14 %Cpu1  : 85.4 us, 14.6 sy,  0.0 ni,  0.0 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st

```

```

15 MiB Mem :   3924.0 total,   1197.7 free,    587.1 used,   2139.2 buff/cache
16 MiB Swap:   3925.0 total,   3924.7 free,     0.3 used.   3044.8 avail Mem
17
18      PID USER      PR  NI    VIRT    RES    SHR S  %CPU  %MEM     TIME+ COMMAND
19    32574 root        20   0    6908    2496    672 R   25.0   0.1   0:01.65 stress-ng-cpu
20    32576 root        20   0    6908    2496    672 R   25.0   0.1   0:01.66 stress-ng-cpu
21    32578 root        20   0   268408 150288    540 R   25.0   3.7   0:01.65 stress-ng-vm
22    32579 root        20   0   268408 152136    540 R   25.0   3.8   0:01.65 stress-ng-vm

```

2.5. 对比两个容器cpu限制情况

```

1  #c1
2  root@castillo:~# docker run -it --rm --name jinliu-c1 --cpus 1 --cpu-shares 800
   lorel/docker-stress-ng --cpu 1 --vm 2
3  stress-ng: info: [1] defaulting to a 86400 second run per stressor
4  stress-ng: info: [1] dispatching hogs: 1 cpu, 2 vm
5
6  root@castillo:~# docker stats
7  CONTAINER ID   NAME                CPU %       MEM USAGE / LIMIT   MEM %       NET I/O
8  2330728e0b6e   sleepy_johnson      0.00%       15.47MiB / 3.832GiB  0.39%       11MB /
   110kB   0B / 121MB   2
9  5640562a51c1   jinliu-c1           97.83%       517.8MiB / 3.832GiB  13.20%       796B / 0B
   0B / 0B   6
10
11 root@castillo:~# top
12 top - 17:13:30 up 10 days,  2:46, 10 users,  load average: 1.23, 1.55, 1.34
13 Tasks: 147 total,  5 running, 142 sleeping,  0 stopped,  0 zombie
14 %Cpu(s): 51.0 us,  0.5 sy,  0.0 ni, 48.5 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
15 MiB Mem :   3924.0 total,   997.4 free,    787.3 used,   2139.3 buff/cache
16 MiB Swap:   3925.0 total,   3924.7 free,     0.3 used.   2844.5 avail Mem
17
18      PID USER      PR  NI    VIRT    RES    SHR S  %CPU  %MEM     TIME+ COMMAND
19    32745 root        20   0   268408 262544    596 R   49.5   6.5   0:37.75 stress-ng-vm
20    32746 root        20   0   268408 262544    596 R   25.2   6.5   0:18.99 stress-ng-vm
21    32742 root        20   0    6908    2432    604 R   24.9   0.1   0:19.01 stress-ng-cpu

```

```
22      2596 root          20    0 1356280  58640  35812 S    0.3   1.5   1:34.86 containerd
```

```
23      10621 castillo    20    0   17300   8072   5660 S    0.3   0.2   0:01.02 sshd
```

```
24
```

```
25 #C2
```

```
26 root@castillo:~# docker run -it --rm --name jinliu-c2 --cpus 1 --cpu-shares 400  
    lorel/docker-stress-ng --cpu 1 --vm 2
```

```
27 stress-ng: info: [1] defaulting to a 86400 second run per stressor
```

```
28 stress-ng: info: [1] dispatching hogs: 1 cpu, 2 vm
```

```
29
```

```
30 root@castillo:~# docker stats
```

```
31 2330728e0b6e    sleepy_johnson    0.00%      15.47MiB / 3.832GiB    0.39%      11MB /  
    110kB    0B / 121MB    2
```

```
32 5640562a51c1    jinliu-c1         100.22%     517.8MiB / 3.832GiB    13.20%     866B / 0B  
    0B / 0B      6
```

```
33 03dd3139eb2d    jinliu-c2         98.85%     314.3MiB / 3.832GiB    8.01%      656B / 0B  
    0B / 0B      6
```

```
34
```

```
35 root@castillo:~# top1
```

```
36 top - 17:15:15 up 10 days,  2:47, 10 users,  load average: 3.74, 2.09, 1.54
```

```
37 Tasks: 153 total,   7 running, 146 sleeping,   0 stopped,   0 zombie
```

```
38 %Cpu(s): 99.0 us,  0.7 sy,  0.0 ni,  0.3 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
```

```
39 MiB Mem :  3924.0 total,   466.8 free,  1315.5 used,  2141.7 buff/cache
```

```
40 MiB Swap:  3925.0 total,  3924.7 free,    0.3 used.  2314.1 avail Mem
```

```
41
```

```
42      PID USER          PR  NI     VIRT    RES    SHR S  %CPU  %MEM    TIME+ COMMAND
```

```
43      32831 root          20   0   268408  262528    580 R   39.3   6.5   0:15.76 stress-ng-vm
```

```
44      32742 root          20   0     6908    2432    604 R   37.7   0.1   0:50.53 stress-ng-cpu
```

```
45      32746 root          20   0   268408  262544    596 R   32.3   6.5   0:48.67 stress-ng-vm
```

```
46      32745 root          20   0   268408  262544    596 R   30.0   6.5   1:21.84 stress-ng-vm
```

```
47      32828 root          20   0     6908    2456    628 R   29.3   0.1   0:12.53 stress-ng-cpu
```

```
48      32832 root          20   0   268408  262528    580 R   29.0   6.5   0:12.90 stress-ng-vm
```

```
49      4154 root          20   0  1595372  84380  50980 S    0.3   2.1   1:29.97 dockerd
```

```
50      31232 root          20   0   712656  11152   8296 S    0.3   0.3   0:03.21 containerd-
```

2.6. systemd验证

```

1 root@castillo:~# docker ps
2 CONTAINER ID   IMAGE                                COMMAND                                  CREATED
3 5565cf04c53b   lorel/docker-stress-ng             "/usr/bin/stress-ng ..."            38
   seconds ago   Up 36 seconds                        jinliu-cl
4 2330728e0b6e   harbor.jinliu.net/myserver/nginx:v1 "nginx -g 'daemon of...'              2 hours
   ago          Up 2 hours                          0.0.0.0:80->80/tcp, :::80->80/tcp, 443/tcp sleepy_johnson
5 root@castillo:~# ps -ef |grep nginx
6 root          31251      31232    0 15:31 pts/0    00:00:00 nginx: master process nginx -g
   daemon off;
7 nobody        31505      31251    0 15:39 pts/0    00:00:00 nginx: worker process
8 root          33005      32653    0 17:21 pts/8    00:00:00 grep --color=auto nginx
9 root@castillo:~# cat /proc/31251/cpuset
10 /system.slice/docker-
   2330728e0b6e9110151f74b714365945522e66e8c03d723f875e09c112748d6c.scope
11 root@castillo:~# cat /sys/fs/cgroup/system.slice/docker-
   2330728e0b6e9110151f74b714365945522e66e8c03d723f875e09c112748d6c.scope/cpu.max
12 max 100000
13 root@castillo:~# cat /sys/fs/cgroup/system.slice/docker-
   2330728e0b6e9110151f74b714365945522e66e8c03d723f875e09c112748d6c.scope/memory.max
14 max
15 root@castillo:~#

```

3. 部署http协议的harbor镜像仓库

3.1. 优点

- 基于角色的访问控制：用户与Docker镜像仓库通过“项目”进行项目管理，可以对不同的账户设置不同的权限，以实现权限的精细管控。
- 镜像复制：镜像可以在多个Registry实例中复制（同步），可以实现高性能、高可用的镜像服务。
- 图形化用户界面：用户可以通过浏览器来浏览，管理当前Docker镜像仓库，管理项目和镜像等。
- AD/LDAP 支：Harbor可以集成企业内部已有的AD/LDAP，用于鉴权认证管理。
- 审计管理：所有针对镜像仓库的操作都可以被记录追溯，用于审计管理。
- 国际化：已拥有英文、中文、德文、日文和俄文等多语言支持版本。
- RESTful API：提供给管理员对于Harbor更多的操控，使得与其它管理软件集成变得更容易。
- 部署简单：提供在线和离线两种安装工具，也可以安装到vSphere平台(OVA方式)虚拟设备。

3.2. 安装harbor


```
1  #修改域名
2  root@castillo:/apps# ls -l
3  total 739540
4  -rw-rw-r-- 1 castillo castillo 757278514 Oct 30 14:48 harbor-offline-installer-
   v2.6.1.tgz
5  drwxr-xr-x 2 root      root          4096 Oct 30 14:30 nginx
6  root@castillo:/apps# tar -zxf harbor-offline-installer-v2.6.1.tgz
7  root@castillo:/apps# cd harbor
8  root@castillo:/apps/harbor# vim harbor.yml
9  root@castillo:/apps/harbor# grep -A 10 "jinliu" harbor.yml
10 hostname: harbor.jinliu.net
11
12 # http related config
13 http:
14   # port for http, default is 80. If https enabled, this port will redirect to https
   port
15   port: 80
16
17 # https related config
18 https:
19   # https port for harbor, default is 443
   port: 443
20
21 root@castillo:/apps/harbor#
22
23 #配置本地域名解析
24 root@castillo:/apps/harbor# cat /etc/hosts
25 127.0.0.1 localhost
26 127.0.1.1 castillo
27
28 # The following lines are desirable for IPv6 capable hosts
29 ::1      ip6-localhost ip6-loopback
30 fe00::0  ip6-localnet
31 ff00::0  ip6-mcastprefix
32 ff02::1  ip6-allnodes
33 ff02::2  ip6-allrouters
34
35 192.168.31.113 harbor.jinliu.net
36 root@castillo:/apps/harbor#
37
38 #配置域名信任
39 root@castillo:/apps/harbor# cat /etc/docker/daemon.json
40 {
41   "registry-mirrors": ["https://rlyvli78.mirror.aliyuncs.com"],
42   "insecure-registries": ["harbor.jinliu.net", "192.168.31.113"]
43 }
44 root@castillo:/apps/harbor#
45
46 #安装
47 root@castillo:/apps/harbor# bash install.sh
48
49 #登录
50 root@castillo:/apps/harbor# docker login harbor.jinliu.net
51 Username: admin
```

```
52 Password:
53 WARNING! Your password will be stored unencrypted in /root/.docker/config.json.
54 Configure a credential helper to remove this warning. See
55 https://docs.docker.com/engine/reference/commandline/login/#credentials-store
56
57 Login Succeeded
58 root@castillo:/apps/harbor#
```

3.3. 推进镜像

```
1 root@castillo:/apps/harbor# docker tag lorel/docker-stress-ng:latest
  harbor.jinliu.net/myserver/docker-stress-ng:latest
2 root@castillo:/apps/harbor# docker push harbor.jinliu.net/myserver/docker-stress-
  ng:latest
3 The push refers to repository [harbor.jinliu.net/myserver/docker-stress-ng]
4 5f70bf18a086: Pushed
5 ea580b0285fe: Pushed
6 6102f0d2ad33: Pushed
7 latest: digest:
  sha256:5768a5d8e196be8c8fabdda04d937aabe1407f397b2f12e1fea2573e4b9d9bef size: 1563
8 root@castillo:/apps/harbor#
9 root@castillo:/apps/harbor#
10 root@castillo:/apps/harbor#
11 root@castillo:/apps/harbor#
```

